

In [1]: %matplotlib inline

```
import os
import sys
import json

import numpy as np
import pandas as pd

from matplotlib import pyplot as plt
from IPython.display import display
import seaborn as sns
sns.set()

# Add our local functions to the path
sys.path.append(os.path.join(os.getcwd(), 'src'))
from models import evaluation
from data.load_data import (get_country_filepaths,
                             split_features_labels_weights,
                             load_data)
from features import process_features
from features.process_features import get_vif, standardize
from features.process_features import RIAU_BASIC_FEATURES

ALGORITHM_NAME = 'rf'
COUNTRY = 'riau'
TRAIN_PATH, TEST_PATH = get_country_filepaths(COUNTRY)
```

In [2]: # Load training data
X_train, y_train, w_train = split_features_labels_weights(TRAIN_PATH)

```
# summarize Loaded data
print('Data has {:,} rows and {:,} columns' \
      .format(*X_train.shape))

print('Percent poor: {:.1%} \tPercent non-poor: {:.1%}' \
      .format(*y_train.miskin.value_counts(normalize=True, ascending=True)))

# print first 5 rows of data
X_train.head()
```

Data has 6,136 rows and 506 columns
Percent poor: 4.9% Percent non-poor: 95.1%

Out[2]:

	r105	r1701	r1702	r1703	r1704	r1705	r1706	r1707	r1708	r1801	...	kons_307	kons_308	kons_309	kons_310	der_nchild10under	der_nmalesover10	der_
0	1	5	5	5	5	5	5	5	5	0	...	0	0	0	0	0	0	
1	2	5	5	5	5	5	5	5	5	1	...	0	0	0	0	0	3	
2	1	5	5	5	5	5	5	5	5	1	...	0	0	0	0	0	3	
3	2	5	5	5	5	5	5	5	5	1	...	0	0	0	0	1	3	
4	2	5	5	5	5	5	5	5	5	1	...	0	0	0	0	0	1	

5 rows × 506 columns

Random forest with simple features

We'll start with the simple set of features stored in RIAU_BASIC_FEATURES in the process_features module.

In [3]: # Select the basic features we've used previously
selected_columns = RIAU_BASIC_FEATURES
print("X shape with selected columns:", X_train[selected_columns].shape)

X shape with selected columns: (6136, 10)

In [4]: # Create DataFrame of feature importances
def get_feat_imp_df(featimps, index=None, sort=True):
 featimps = pd.DataFrame(featimps, columns=['importance'])
 if index is not None:
 featimps.index = index
 if sort:
 featimps = featimps.sort_values('importance', ascending=False)
 return featimps

```
In [5]: from sklearn.ensemble import RandomForestClassifier

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH,
                                      selected_columns=selected_columns)

# Fit the model
model = RandomForestClassifier()
%time model.fit(X_train, y_train)

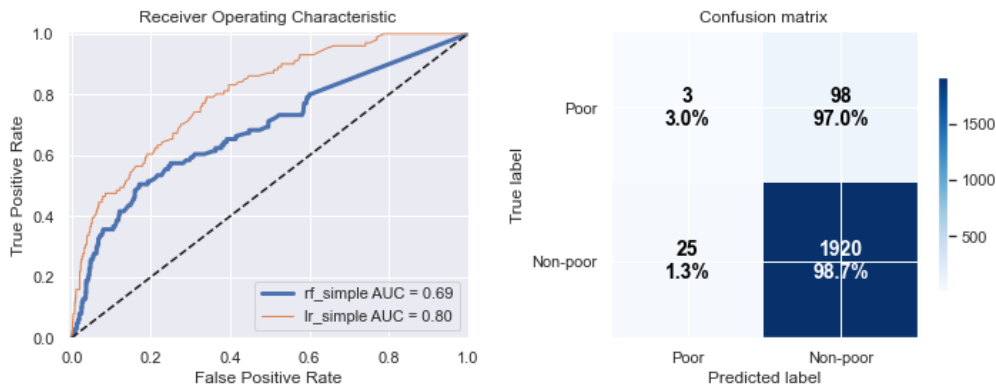
# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.2%}".format(score))
featimps = get_feat_imp_df(model.feature_importances_, index=X_train.columns)

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH,
                                   selected_columns=selected_columns)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                   compare_models=['lr_simple'],
                                   store_model=True,
                                   model_name='simple',
                                   prefix=ALGORITHM_NAME,
                                   country=COUNTRY,
                                   model=model,
                                   features=featimps)
```

Wall time: 330 ms
Wall time: 91 ms
In-sample score: 96.99%



Model Scores

	rf_simple	lr_simple
accuracy	0.939883	0.950147
recall	0.029703	0.009901
precision	0.107143	0.333333
f1	0.046512	0.019231
cross_entropy	0.492030	0.164615
roc_auc	0.685294	0.799743
cohen_kappa	0.025631	0.016430

Actual poverty rate: 6.82%
Predicted poverty rate: 3.16%

In [6]:

featimps

Out[6]:

	importance
der_nbekerja	0.175885
der_nliterate	0.161012
r1816	0.132766
r301	0.132566
der_nchild10under	0.115980
der_nmalesover10	0.093693
r1809a	0.086772
der_nfemalesover10	0.084614
kons_305	0.008983
kons_262	0.007727

Random forest with all features

Now let's apply the Random Forest classifier to the dataset with all features.

```
In [7]: # Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)

# Fit the model
model = RandomForestClassifier()
%time model.fit(X_train, y_train)

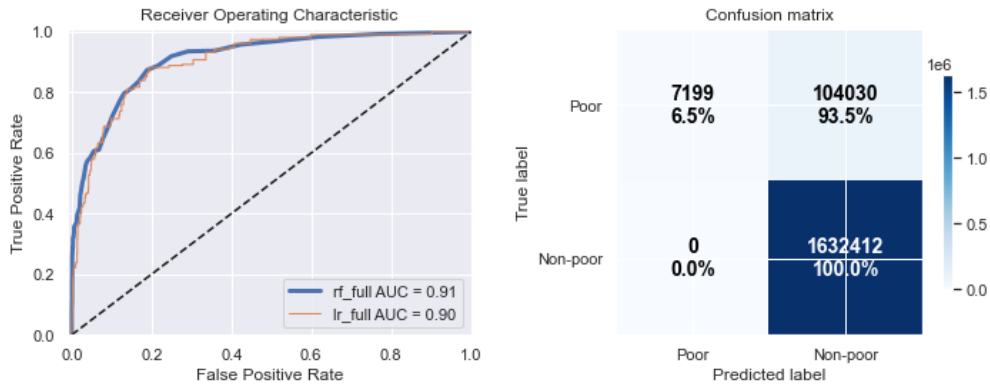
# Get an initial score
%time score = model.score(X_train, y_train, w_train)
print("In-sample score: {:.2%}".format(score))
featimps = get_feat_imp_df(model.feature_importances_, index=X_train.columns)

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob, w_test,
                                   compare_models='lr_full',
                                   store_model=True,
                                   model_name='full',
                                   prefix=ALGORITHM_NAME,
                                   country=COUNTRY,
                                   model=model,
                                   features=featimps)
```

Wall time: 1.23 s
Wall time: 111 ms
In-sample score: 100.00%

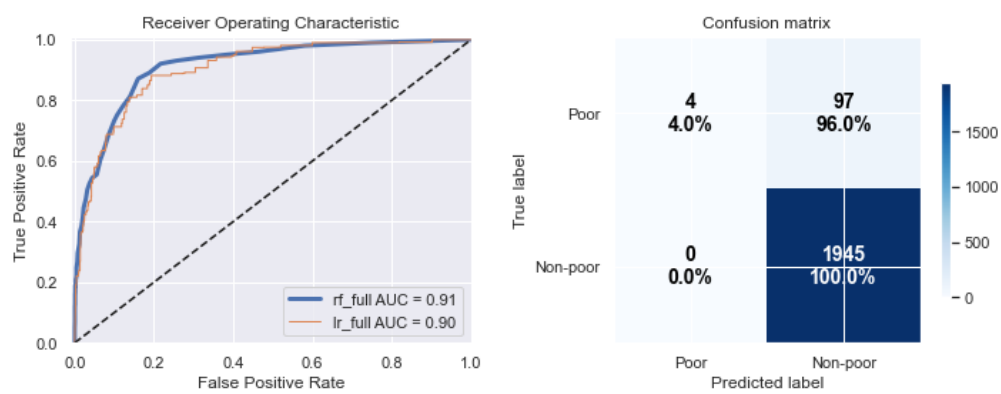


Model Scores

	rf_full	lr_full
accuracy	0.940338	0.930738
recall	0.064722	0.466875
precision	1.000000	0.457938
f1	0.121576	0.462363
cross_entropy	0.167726	1.327124
roc_auc	0.912237	0.904081
cohen_kappa	0.072703	0.363980

Actual poverty rate: 6.82%
Predicted poverty rate: 5.35%

```
In [8]: # Evaluate performance and store model without w_test
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                   compare_models='lr_full',
                                   store_model=True,
                                   model_name='full',
                                   prefix=ALGORITHM_NAME,
                                   country=COUNTRY,
                                   model=model,
                                   features=featimps)
```



Model Scores

```
In [9]: featimps
```

Out[9]:

	importance
r301	0.029641
der_nchild10under	0.023134
r303	0.019257
r1804	0.018273
der_ninternetpast3mo	0.015831
...	...
kons_15	0.000000
r1901c	0.000000
kons_54	0.000000
kons_218	0.000000
kons_7	0.000000

506 rows x 1 columns

Class Weighting

The random forest classifier includes the option to balance the classes with `class_weight='balanced'`, so we'll first try this before applying a sampling method. We'll also pick a few parameters to reduce overfitting.

```
In [10]: # Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)

# Fit the model
model = RandomForestClassifier(n_estimators=100,
                              max_depth=20,
                              min_samples_leaf=5,
                              min_samples_split=5,
                              class_weight='balanced')
%time model.fit(X_train, y_train)

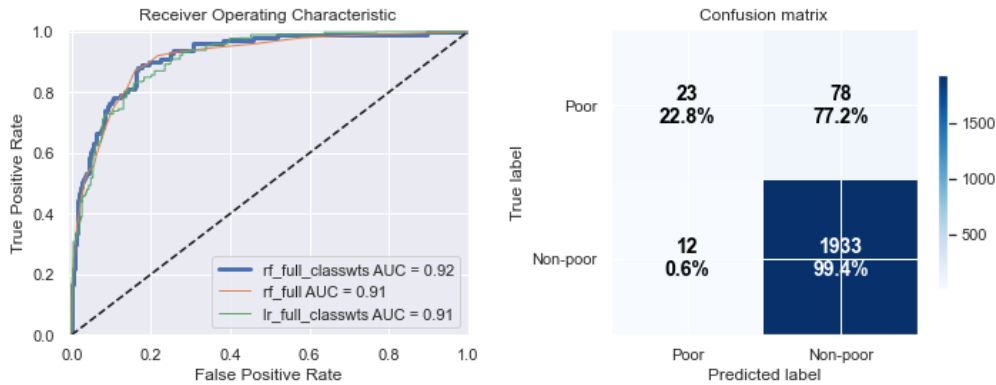
# Get an initial score
%time score = model.score(X_train, y_train, w_train)
print("In-sample score: {:.2%}".format(score))
featimps = get_feat_imp_df(model.feature_importances_, index=X_train.columns)

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                   compare_models=['rf_full',
                                                  'lr_full_classwts'],
                                   store_model=True,
                                   model_name='full_classwts',
                                   prefix=ALGORITHM_NAME,
                                   country=COUNTRY,
                                   model=model,
                                   features=featimps)
```

Wall time: 1.17 s
Wall time: 129 ms
In-sample score: 99.71%



Model Scores

	rf_full_classwts	rf_full	lr_full_classwts
accuracy	0.956012	0.952590	0.929499
recall	0.227723	0.039604	0.479928
precision	0.657143	1.000000	0.450620
f1	0.338235	0.076190	0.464812
cross_entropy	0.172160	0.145552	1.082523
roc_auc	0.921143	0.913816	0.914168
cohen_kappa	0.320983	0.072703	0.400154

Actual poverty rate: 6.82%
Predicted poverty rate: 6.07%

Undersampling

Now we'll apply undersampling and see how the performance is affected.

```
In [12]: from imblearn.under_sampling import RandomUnderSampler

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)
cols = X_train.columns

# Apply random undersampling
X_train, y_train = RandomUnderSampler().fit_resample(X_train, y_train)
X_train = pd.DataFrame(X_train, columns=cols)
print("X shape after undersampling: ", X_train.shape)

# Fit the model
model = RandomForestClassifier(n_estimators=100,
                              max_depth=20,
                              min_samples_leaf=5,
                              min_samples_split=5)

%time model.fit(X_train, y_train)


# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.0.2%}".format(score))
featimps = get_feat_imp_df(model.feature_importances_, index=X_train.columns)

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
%time y_pred = model.predict(X_test)
%time y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                    compare_models=['rf_full_classwts',
                                                    'lr_full_oversample_cv'],
                                    store_model=True,
                                    model_name='full_undersample',
                                    prefix=ALGORITHM_NAME,
                                    country=COUNTRY,
                                    model=model,
                                    features=featimps)
```

X shape after undersampling: (602, 506)
Wall time: 208 ms
Wall time: 30 ms
In-sample score: 96.18%
Wall time: 53 ms
Wall time: 40 ms



Model Scores

	rf_full_undersample	rf_full_classwts	lr_full_oversample_cv
accuracy	0.801075	0.956012	0.940535
recall	0.871287	0.227723	0.682982
precision	0.182573	0.657143	0.526119
f1	0.301887	0.338235	0.594375
cross_entropy	0.449701	0.172160	0.282615
roc_auc	0.913039	0.921143	0.932446
cohen_kappa	0.239838	0.320983	0.501967

Actual poverty rate: 6.82%
Predicted poverty rate: 25.53%

Oversampling

```
In [13]: from imblearn.over_sampling import SMOTE

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)
cols = X_train.columns

# Apply oversampling with SMOTE
X_train, y_train = SMOTE().fit_resample(X_train, y_train)
X_train = pd.DataFrame(X_train, columns=cols)
print("X shape after oversampling: ", X_train.shape)

# Fit the model
model = RandomForestClassifier(n_estimators=100,
                              max_depth=20,
                              min_samples_leaf=5,
                              min_samples_split=5)
%time model.fit(X_train, y_train)

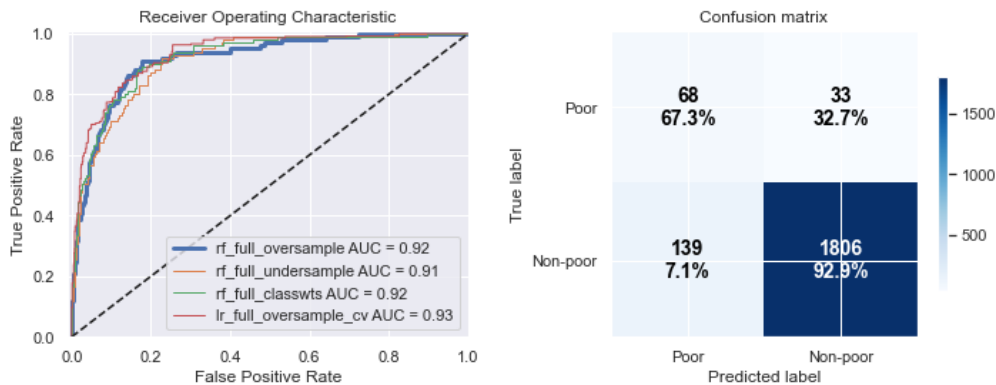
# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.2%}".format(score))

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
%time y_pred = model.predict(X_test)
%time y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                    compare_models=['rf_full_undersample',
                                                    'rf_full_classwts',
                                                    'lr_full_oversample_cv'],
                                    store_model=True,
                                    model_name='full_oversample',
                                    prefix=ALGORITHM_NAME,
                                    country=COUNTRY,
                                    model=model,
                                    features=featimps)
```

X shape after oversampling: (11670, 506)
Wall time: 3.85 s
Wall time: 216 ms
In-sample score: 99.02%
Wall time: 59 ms
Wall time: 44 ms



Model Scores

	rf_full_oversample	rf_full_undersample	rf_full_classwts	lr_full_oversample_cv
accuracy	0.915934	0.801075	0.956012	0.940535
recall	0.673267	0.871287	0.227723	0.682982
precision	0.328502	0.182573	0.657143	0.526119
f1	0.441558	0.301887	0.338235	0.594375
cross_entropy	0.378935	0.449701	0.172160	0.282615
roc_auc	0.917188	0.913039	0.921143	0.932446
cohen_kappa	0.401870	0.239838	0.320983	0.501967

Actual poverty rate: 6.82%
Predicted poverty rate: 6.47%

Tune parameters and cross-validate

```
In [14]: from sklearn.model_selection import GridSearchCV

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)
cols = X_train.columns

# Apply random undersampling
X_train, y_train = RandomUnderSampler().fit_resample(X_train, y_train)
X_train = pd.DataFrame(X_train, columns=cols)
print("X shape after undersampling: ", X_train.shape)

# build the model
estimator = RandomForestClassifier()
parameters = {'n_estimators': [10, 50, 100],
              'max_depth': np.arange(1,16,5),
              'min_samples_split': np.arange(2,21,10),
              'min_samples_leaf': np.arange(1,46,20)}
model = GridSearchCV(estimator, parameters, verbose=1, cv=5, n_jobs=-1)
%time model.fit(X_train, y_train)

# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.2%}".format(score))
print("Best model parameters:", model.best_params_)
featimps = get_feat_imp_df(model.best_estimator_.feature_importances_, index=X_train.columns)

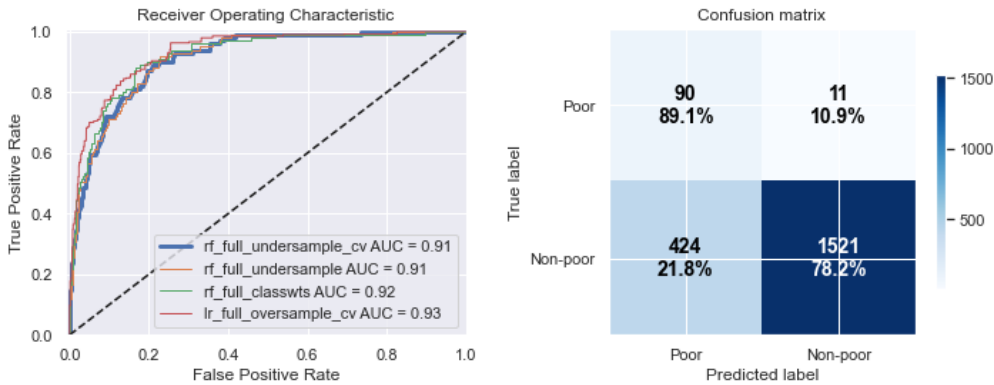
# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                   compare_models=['rf_full_undersample',
                                                  'rf_full_classwts',
                                                  'lr_full_oversample_cv'],
                                   store_model=True,
                                   model_name='full_undersample_cv',
                                   prefix=ALGORITHM_NAME,
                                   country=COUNTRY,
                                   model=model,
                                   features=featimps)

best_model = model.best_estimator_
```

X shape after undersampling: (602, 506)
Fitting 5 folds for each of 54 candidates, totalling 270 fits
Wall time: 9.24 s
Wall time: 26 ms
In-sample score: 100.00%
Best model parameters: {'max_depth': 11, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}



Model Scores

	rf_full_undersample_cv	rf_full_undersample	rf_full_classwts	lr_full_oversample_cv
accuracy	0.787390	0.801075	0.956012	0.940535
recall	0.891089	0.871287	0.227723	0.682982
precision	0.175097	0.182573	0.657143	0.526119
f1	0.292683	0.301887	0.338235	0.594375
cross_entropy	0.447345	0.449701	0.172160	0.282615
roc_auc	0.913467	0.913039	0.921143	0.932446
cohen_kappa	0.229069	0.239838	0.320983	0.501967

Actual poverty rate: 6.82%
Predicted poverty rate: 26.71%

Cross validation for this model has only a minor impact on performance. As with Logistic Regression, we tend to get better recall using undersampling, but better overall performance using oversampling.

In [15]:

featimps.head(20)

Out[15]:

	importance
r301	0.043973
der_nchild10under	0.034253
r303	0.024580
der_nliterate	0.019845
r304	0.017684
r2002	0.017249
kons_213	0.016827
r1808	0.016146
r1701	0.015258
r2001b	0.015079
kons_107	0.014842
r1804	0.014748
kons_166	0.014543
kons_268	0.014061
kons_214	0.013230
der_ninternetpast3mo	0.013063
kons_229	0.011620
r1809d	0.011051
der_nfemalesover10	0.010179
kons_56	0.009068

AdaBoost Random Forest Classifier

```
In [16]: from sklearn.ensemble import AdaBoostClassifier

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH)
cols = X_train.columns

# Apply random undersampling
X_train, y_train = RandomUnderSampler().fit_resample(X_train, y_train)
X_train = pd.DataFrame(X_train, columns=cols)
print("X shape after undersampling: ", X_train.shape)

# build the model
estimator = AdaBoostClassifier(best_model)

parameters = {'n_estimators': [50, 100],
              'learning_rate': [0.01, 0.1]}
model = GridSearchCV(estimator, parameters, verbose=3, cv=3, n_jobs=-1)
%time model.fit(X_train, y_train)

# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.2%}".format(score))
print("Best model parameters:", model.best_params_)
featimps = get_feat_imp_df(model.best_estimator_.feature_importances_, index=X_train.columns)


# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                    compare_models=['rf_full_undersample_cv',
                                                    'rf_full_undersample',
                                                    'rf_full_classwts',
                                                    'lr_full_oversample_cv'],
                                    store_model=True,
                                    model_name='full_undersample_cv_ada',
                                    prefix=ALGORITHM_NAME,
                                    country=COUNTRY,
                                    model=model,
                                    features=featimps)

best_model = model.best_estimator_
```

X shape after undersampling: (602, 506)
Fitting 3 folds for each of 4 candidates, totalling 12 fits
Wall time: 1.23 s
Wall time: 23 ms
In-sample score: 100.00%
Best model parameters: {'learning_rate': 0.01, 'n_estimators': 100}



Model Scores

	rf_full_undersample_cv_ada	rf_full_undersample_cv	rf_full_undersample	rf_full_classwts	lr_full_oversample_cv
accuracy	0.809873	0.787390	0.801075	0.956012	0.940535
recall	0.821782	0.891089	0.871287	0.227723	0.682982
precision	0.182819	0.175097	0.182573	0.657143	0.526119
f1	0.299099	0.292683	0.301887	0.338235	0.594375
cross_entropy	0.432333	0.447345	0.449701	0.172160	0.282615
roc_auc	0.913991	0.913467	0.913039	0.921143	0.932446
cohen_kappa	0.237519	0.229069	0.239838	0.320983	0.501967

Actual poverty rate: 6.82%
Predicted poverty rate: 24.64%

In [17]:

featimps.head(20)

Out[17]:

	importance
r301	0.042055
der_nchild10under	0.028776
r303	0.020379
r1804	0.019942
kons_166	0.016315
r1701	0.016169
der_ninternetpast3mo	0.015275
r1702	0.014465
der_nliterate	0.014016
kons_214	0.013646
kons_107	0.013027
kons_56	0.012429
r2001f	0.012314
r1803	0.012229
r1703	0.011794
r1808	0.011514
kons_229	0.011435
r2001k	0.010665
kons_236	0.010239
der_nfemalesover10	0.009768

Feature selection

```
In [18]: feats = featimps[featimps.cumsum() <= 0.80].dropna().index.values

# Load and transform the training data
X_train, y_train, w_train = load_data(TRAIN_PATH, selected_columns=feats)
cols = X_train.columns
print("X shape after feature selection: ", X_train.shape)

# Apply random undersampling
X_train, y_train = RandomUnderSampler().fit_resample(X_train, y_train)
X_train = pd.DataFrame(X_train, columns=cols)
print("X shape after undersampling: ", X_train.shape)

# Fit the model
model = best_model
%time model.fit(X_train, y_train)


# Get an initial score
%time score = model.score(X_train, y_train)
print("In-sample score: {:.2%}".format(score))
featimps = get_feat_imp_df(model.feature_importances_, index=X_train.columns)

# Load the test set
X_test, y_test, w_test = load_data(TEST_PATH, selected_columns=feats)

# Run the model
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:,:1]

# Evaluate performance and store model
metrics = evaluation.evaluate_model(y_test, y_pred, y_prob,
                                    compare_models=[ 'rf_full_undersample_cv_ada',
                                                      'rf_full_undersample_cv',
                                                      'rf_full_undersample',
                                                      'rf_full_classwts',
                                                      'lr_full_oversample_cv'],
                                    store_model=True,
                                    model_name='feats',
                                    prefix=ALGORITHM_NAME,
                                    country=COUNTRY,
                                    model=model,
                                    features=featimps)
```

X shape after feature selection: (6136, 168)
X shape after undersampling: (602, 168)
Wall time: 407 ms
Wall time: 39 ms
In-sample score: 100.00%



Model Scores

	rf_feats	rf_full_undersample_cv_ada	rf_full_undersample_cv	rf_full_undersample	rf_full_classwts	lr_full_oversample_cv
accuracy	0.804008	0.809873	0.787390	0.801075	0.956012	0.940535
recall	0.881188	0.821782	0.891089	0.871287	0.227723	0.682982
precision	0.186192	0.182819	0.175097	0.182573	0.657143	0.526119
f1	0.307427	0.299099	0.292683	0.301887	0.338235	0.594375
cross_entropy	0.425581	0.432333	0.447345	0.449701	0.172160	0.282615
roc_auc	0.919494	0.913991	0.913467	0.913039	0.921143	0.932446
cohen_kappa	0.245968	0.237519	0.229069	0.239838	0.320983	0.501967

Actual poverty rate: 6.82%
Predicted poverty rate: 24.95%

```
In [19]: cons_feats = [x for x in feats if x[0:5] != 'kons_']  
print("{} consumables features selected:".format(len(cons_feats)))  
for x in cons_feats:  
    print(x)
```


106 consumables features selected:

kons_166
kons_214
kons_107
kons_56
kons_229
kons_236
kons_268
kons_81
kons_277
kons_179
kons_283
kons_207
kons_278
kons_213
kons_224
kons_273
kons_91
kons_234
kons_78
kons_180
kons_34
kons_259
kons_160
kons_148
kons_264
kons_48
kons_162
kons_300
kons_6
kons_68
kons_26
kons_111
kons_282
kons_155
kons_165
kons_195
kons_301
kons_168
kons_138
kons_145
kons_82
kons_96
kons_258
kons_13
kons_131
kons_256
kons_170
kons_299
kons_176
kons_135
kons_95
kons_254
kons_238
kons_74
kons_163
kons_73
kons_129
kons_83
kons_266
kons_287
kons_257
kons_280
kons_216
kons_80
kons_261
kons_18
kons_184
kons_156
kons_103
kons_113
kons_157
kons_57
kons_147
kons_42
kons_141
kons_140
kons_243
kons_284
kons_164
kons_79
kons_30
kons_194
kons_253
kons_159
kons_33
kons_144
kons_84
kons_85
kons_136
kons_153
kons_93

kons_117
kons_175
kons_115
kons_154
kons_123
kons_271
kons_10
kons_251
kons_193
kons_76
kons_275
kons_21
kons_75
kons_269
kons_86

Random Forest Riau Summary


```
In [20]: evaluation.compare_algorithm_models(ALGORITHM_NAME, COUNTRY)
```



Model Scores

VS Logistic Regression

```
In [21]: evaluation.compare_algorithm_models('lr', COUNTRY)
```



Model Scores

	accuracy	recall	precision	f1	cross_entropy	roc_auc	cohen_kappa	pov_rate_error
lr_full	0.930738	0.466875	0.457938	0.462363	1.327124	0.904081	0.363980	0.000409
lr_full_classwts	0.929499	0.479928	0.450620	0.464812	1.082523	0.914168	0.400154	0.004715
lr_full_oversample	0.933880	0.632532	0.485975	0.549652	0.364062	0.927669	0.452221	0.018465
lr_full_oversample_cv	0.940535	0.682982	0.526119	0.594375	0.282615	0.932446	0.501967	0.021896
lr_full_undersample	0.835955	0.889434	0.265465	0.408891	0.634811	0.933625	0.320408	0.154480
lr_l1_feats_oversample_cv	0.854298	0.823472	0.280952	0.418963	0.364059	0.915807	0.346496	0.124128
lr_simple_classwts	0.732649	0.702970	0.120748	0.206096	0.537430	0.804167	0.133050	0.341761